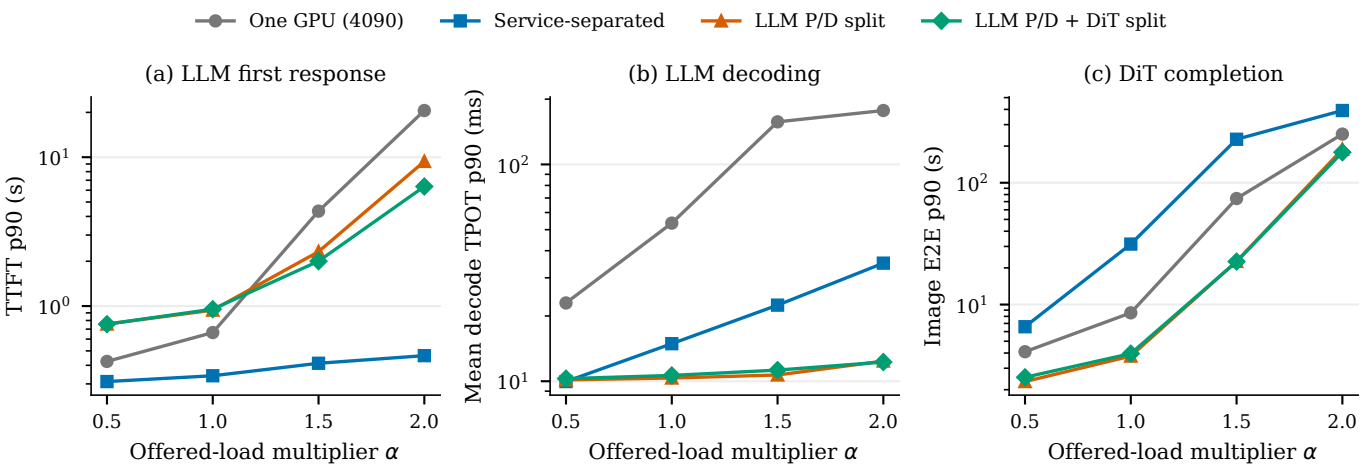
One GPU (4090)
Service-separated
LLM P/D split
LLM P/D + DiT split
(a) LLM first response
TTFT p90 (s)
10^1
10^0
0.5
1.0
1.5
2.0
Offered-load multiplier α
(b) LLM decoding
Mean decode TPOT p90 (ms)
10^2
10^1
0.5
1.0
1.5
2.0
Offered-load multiplier α
(c) DiT completion
Image E2E p90 (s)
10^2
10^1
0.5
1.0
1.5
2.0
Offered-load multiplier α